

Factor Extraction from Macroeconomic News Streams to Drive Agentic Financial Trading Strategies

Chuan Bin Phoe, Neaton Ang

Problem Statement

Motivation

- Traditional time-series models tend to ignore qualitative macro and industry narratives
- Financial news is noisy, unstructured, and difficult to translate into actionable signals
- We need a structured pipeline that extracts relevant news, summarizes key points, and generates interpretable sentiment for forecasting

Training - Data Loader

OOP: TrainDataLoader class

```
config = Config("config.env")
print(f"Config: {config}")
train_data_loader = TrainDataLoader(config)
```

Download Bloomberg News
data using HF datasets



Bloomberg financial
news dataset

446k rows



Data Validation using
Pydantic Class



Pydantic objects

Ensure data integrity



Parquet is a binary,
columnar storage format
Cloud ready: Speed and
Compression > CSV

Data Caching in Parquet files

Subsequent requests for that data can
be served more quickly without needing
to retrieve it from the original source

```
Downloading dataset 'danidanou/Bloomberg_Financial_News' with split 'train' from the Hugging Face Hub...
bloomberg_financial_data.parquet.zip: 100% [██████████] 482M/482M [00:02<00:00, 430MB/s]
Generating train split: 100% [██████████] 446762/446762 [00:08<00:00, 51255.07 examples/s]
--- Download Successful! ---
```

```
--- Starting validation of 446762 entries ---
Validating entries: 100% [██████████] 446762/446762 [00:28<00:00, 15910.17it/s]
--- Validation Complete! ---
Training dataset validated.
Saving processed dataset to local cache at '../data/danidanou_Bloomberg_Financial_News_train'...
Total number of rows: 446762
Loading pipeline completed.
```

Dataset	Size on Amazon S3	Query Run Time	Data Scanned	Cost
Data stored as CSV files	1 TB	236 seconds	1.15 TB	\$5.75
Data stored in Apache Parquet Format	130 GB	6.78 seconds	2.51 GB	\$0.01
Savings	87% less when using Parquet	34x faster	99% less data scanned	99.7% savings

Databricks <https://www.databricks.com/glossary/what-is-parquet>

Training - Data Processor

OOP: NewsProcessor class

```
processor = NewsProcessor(config)
```

Processing Pipeline

```
from processor import NewsProcessor
import nest_asyncio
nest_asyncio.apply()

processor = NewsProcessor(config)
sample = processor.remove_redundant_info(train_ds[DATA_START:DATA_END])
df = processor.enrich_news_entries_with_classifications(sample, save_path=f"{GDRIVE_PATH}processed_news")
df = processor.group_by_date_and_industry(df, save_path=f"{GDRIVE_PATH}grouped_news")
df = processor.filter_and_analyze_news(df)
df = processor.extract_impactful_news(df, top_n=3, save_path=f"{GDRIVE_PATH}impact_news")
df = processor.get Consolidated sentiment(df, save_path=f"{GDRIVE_PATH}sentiment_news")
final_df = await processor.get_explanation(df, save_path=f"{GDRIVE_PATH}explanation_news")
```

Auto-saving to Google Drive at every step

Headline	Link	Date	Article		
1					
Headline	Link	Date	Article	Industry	Sentiment

1

Load HF weights and run inference to get:



Hugging Face

Sentiment

ProsusAI / **FinBERT**



Industry

MoritzLaurer /
mDeBERTa-v3-base-mnli-xnli

```
FinBERT Sentiment: 100%|██████████| 1550/1550 [53:04<00:00,  2.05s/batch]
Industry Classification: 100%|██████████| 12399/12399 [6:27:58<00:00,  1.88s/batch]
Completed processing 396762 entries
```

Ran on Google Colab's Nvidia A100 GPU

2

Grouping all articles with same (Date, Industry) into same row
Filtering out Industry=None news, get analysis statistics

Training - Data Processor

OOP: NewsProcessor class

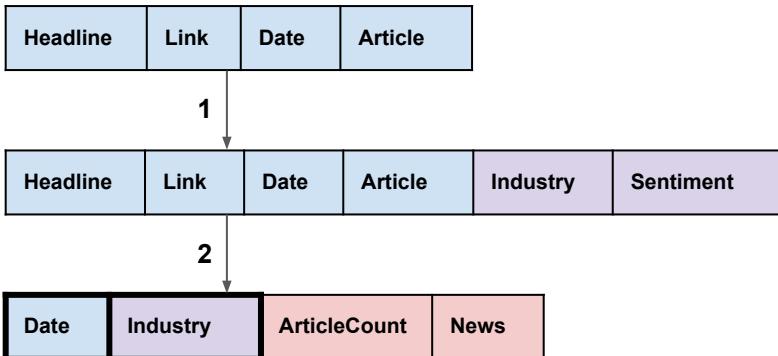
```
processor = NewsProcessor(config)
```

Processing Pipeline

```
from processor import NewsProcessor
import nest_asyncio
nest_asyncio.apply()

processor = NewsProcessor(config)
sample = processor.remove_redundant_info(train_ds[DATA_START:DATA_END])
df = processor.enrich_news_entries_with_classifications(sample, save_path=f"{GDRIVE_PATH}processed_news")
df = processor.group_by_date_and_industry(df, save_path=f"{GDRIVE_PATH}grouped_news")
df = processor.filter_and_analyze_news(df)
df = processor.extract_impactful_news(df, top_n=3, save_path=f"{GDRIVE_PATH}impact_news")
df = processor.get Consolidated sentiment(df, save_path=f"{GDRIVE_PATH}sentiment_news")
final_df = await processor.get_explanation(df, save_path=f"{GDRIVE_PATH}explanation_news")
```

Auto-saving to Google Drive at every step



Each row is 1 unique (Date, Industry) pair, News contains all articles for that pair

1

Load HF weights and run inference to get:



Hugging Face

Sentiment

ProsusAI / **FinBERT**

Google
colab

Industry

MoritzLaurer /
mDeBERTa-v3-base-mnli-xnli

```
FinBERT Sentiment: 100% | 1550/1550 [53:04<00:00, 2.05s/batch]
Industry Classification: 100% | 12399/12399 [6:27:58<00:00, 1.88s/batch]
Completed processing 396762 entries
```

Ran on Google Colab's Nvidia A100 GPU

2

Grouping all articles with same (Date, Industry) into same row
Filtering out Industry=None news, get analysis statistics

```
=====
Dropped 1141 (Industry, Date) pairs with Industry='None'
Remaining pairs: 13591
=====

Summary Statistics:
Total unique (Industry, Date) pairs: 13591
Average articles per pair: 27.91
Max articles in a pair: 1262
Min articles in a pair: 1
25th percentile: 4.0
50th percentile: 13.0
75th percentile: 29.0
Number of pairs with at least 3 articles: 11284
Total articles: 379362
```

Training - Data Processor

OOP: NewsProcessor class

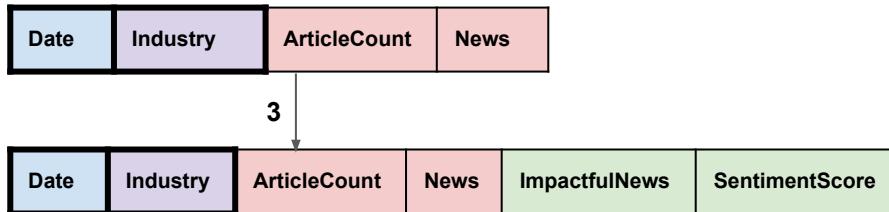
```
processor = NewsProcessor(config)
```

Processing Pipeline

```
from processor import NewsProcessor
import nest_asyncio
nest_asyncio.apply()

processor = NewsProcessor(config)
sample = processor.remove_redundant_info(train_ds[DATA_START:DATA_END])
df = processor.enrich_news_entries_with_classifications(sample, save_path=f"{GDRIVE_PATH}processed_news")
df = processor.group_by_date_and_industry(df, save_path=f"{GDRIVE_PATH}grouped_news")
df = processor.filter_and_analyze_news(df)
df = processor.extract_impactful_news(df, top_n=3, save_path=f"{GDRIVE_PATH}impact_news")
df = processor.get Consolidated sentiment(df, save_path=f"{GDRIVE_PATH}sentiment_news")
final_df = await processor.get_explanation(df, save_path=f"{GDRIVE_PATH}explanation_news")
```

Auto-saving to Google Drive at every step



3

Get 3 most impactful news for each row (based on abs(sentiment score))

Rerun sentiment scoring with FinBERT based on them

- Market psychology: Big headlines move markets
- Cost constraints: Reduce input tokens for next step

```
Extracting top 3 impactful news per (Industry, Date) pair...
Processing groups: 100% [██████] 13591/13591 [00:00<00:00, 15013.76it/s]
Processing 13591 news entries...
FinBERT Sentiment: 100% [██████] 54/54 [01:46<00:00, 1.96s/batch]
Completed processing 13591 entries
```

4

Asynchronous LLM call to generate sentiment explanation

Based on:

- Impactful News
- FinBERT sentiment score
- General Market News from the same day

Output Validation using Pydantic Class

Training - Data Processor

OOP: NewsProcessor class

```
processor = NewsProcessor(config)
```

Processing Pipeline

```
from processor import NewsProcessor
import nest_asyncio
nest_asyncio.apply()

processor = NewsProcessor(config)
sample = processor.remove_redundant_info(train_ds[DATA_START:DATA_END])
df = processor.enrich_news_entries_with_classifications(sample, save_path=f'{GDRIVE_PATH}processed_news')
df = processor.group_by_date_and_industry(df, save_path=f'{GDRIVE_PATH}grouped_news')
df = processor.filter_and_analyze_news(df)
df = processor.extract_impactful_news(df, top_n=3, save_path=f'{GDRIVE_PATH}impact_news')
df = processor.get Consolidated sentiment(df, save_path=f'{GDRIVE_PATH}sentiment_news')
final_df = await processor.get_explanation(df, save_path=f'{GDRIVE_PATH}explanation_news')
```

Auto-saving to Google Drive at every step

Date	Industry	ArticleCount	News	ImpactfulNews	SentimentScore
------	----------	--------------	------	---------------	----------------

4

	Industry	Date	News	ArticleCount	ImpactfulNews	AvgSentimentScore	SentimentScore	SentimentExplanation
0	Communication Services	2011-10-06	[{"Headline": "FCC to Revamp Phone Subsidy to ...	2	[{"Headline": "Euro-Area Leaders to Hold Summit...]	0.709191	-0.291917	Overall sentiment for the Communications Servi...
1	Consumer Discretionary	2011-10-06	[{"Headline": "PepsiCo May Purchase Russian Dr...]	1	[{"Headline": "PepsiCo May Purchase Russian Dr...]	0.881740	0.888237	The article's sentiment is strongly positive f...
2	Consumer Staples	2011-10-06	[{"Headline": "Ukraine's Grain Harvest Advance...]	1	[{"Headline": "Ukraine's Grain Harvest Advance...]	-0.918589	-0.917441	Explanation: FinBERT indicates a strongly nega...
3	Energy	2011-10-06	[{"Headline": "Clean-Tech Companies Should Get...]	9	[{"Headline": "Norway Boosts Mongstad Carbon-S...]	0.252093	-0.252850	The energy-angle sentiment is mildly negative,...

3

Get 3 most impactful news for each row (based on abs(sentiment score))

Rerun sentiment scoring with FinBERT based on them

- Market psychology: Big headlines move markets
- Cost constraints: Reduce input tokens for next step

```
Extracting top 3 impactful news per (Industry, Date) pair...
Processing groups: 100% [██████] 13591/13591 [00:00<00:00, 15013.76it/s]
Processing 13591 news entries...
FinBERT Sentiment: 100% [██████] 54/54 [01:46<00:00, 1.96s/batch]
Completed processing 13591 entries
```

4

Asynchronous LLM call to generate sentiment explanation

Based on:

- Impactful News
- FinBERT sentiment score
- General Market News from the **same day**

Output Validation using Pydantic Class

Forecasting Model

Modelling Objective.

Quantify whether structured news sentiment provides incremental predictive power beyond traditional market & price-based factors.

Data Overview.

- ~7 years of daily stock price data (AAPL, MSFT, AMZN, ... + SPY)
- Daily returns, rolling volatilities, momentum, liquidity, etc.
- Industry assignment - (Agentic LLM Extracted)
- Key-point extraction - (Agentic LLM Extracted)
- Sentiment scoring (per-article + aggregated) - (Agentic LLM Extracted)

Experimental Setup

- **Data split: Walk-Forward split (Quantile at Date)**
 - Chronological split: the first 70% of dates are used for training and the remaining 30% of dates are used for testing
 - Model is trained only on past data and evaluated on a contiguous future block of data
 - No shuffling
 - This ensures that there is no leakage from the test period into training
- **Classification Metrics**
 - **Accuracy:** % of correct predictions; use when classes balanced and all errors equally costly
 - **AUC:** Ranks positives vs negatives across all thresholds; robust to class imbalance
- **Forecasting Metric**
 - **R-Square:** Fraction of target variance explained by predictions; standard for continuous-value forecasting (e.g., returns)

Experiment Setup: Baseline Clarification

- We use a “CAPM-style market-only baseline”, not the theoretical CAPM model
- It is simply a single-factor forecasting model using only forward SPY return
 - “Forward SPY return” = the cumulative percentage move of SPY from today to h days ahead
- This will show how much predictive structure exists without sentiment, factors, or non-linearities

Model V1: Trained on Indexes

We first train return-forecasting models on indexes.

SPY → market proxy

Sector ETFs (e.g., XLF, XLK) → Represents an industry

Using only.

- Market Factor (Forward SPY)
- Momentum/Volatility features
- Sentiment features

Purpose.

Understand baseline signal strength before going to noisier individual stocks

Results.

Market-only baseline (CAPM-style) wins index have betas close to 1 and highly correlated with SPY, so the CAPM market factor captures a large portion of their return variance.

Model V2: Trained on Individual Stocks

We then transition to individual stocks.

Harder problem → Lower autocorrelation in returns, sentiment may only matter on event-heavy days → but it is more realistic

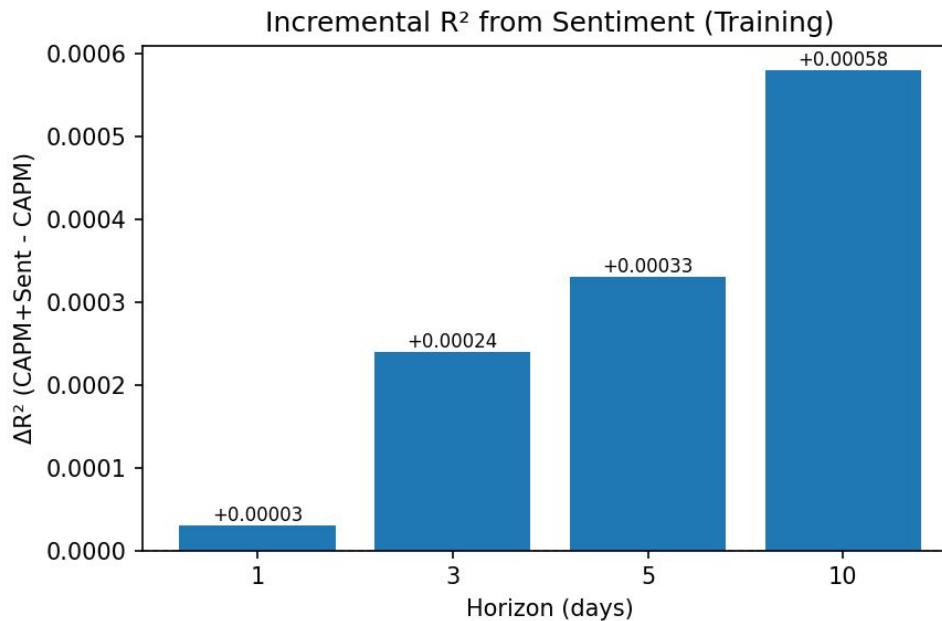
Key engineered features:

Category	Features
Price / Liquidity	ret, ret_3d, ret_5d, ret_vol_20d, ret_vol_60d, dollar_vol_rel_3d, dollar_vol_rel_20d
Sentiment	Standardized daily score, Rolling 3d/5d means & sums, Shock metric = (today – 5d mean)/std, News count z-score
Forward targets	ret_next_1d, ret_next_3d, ret_next_5d, ret_next_10d

Model V2: 4 Main Variants

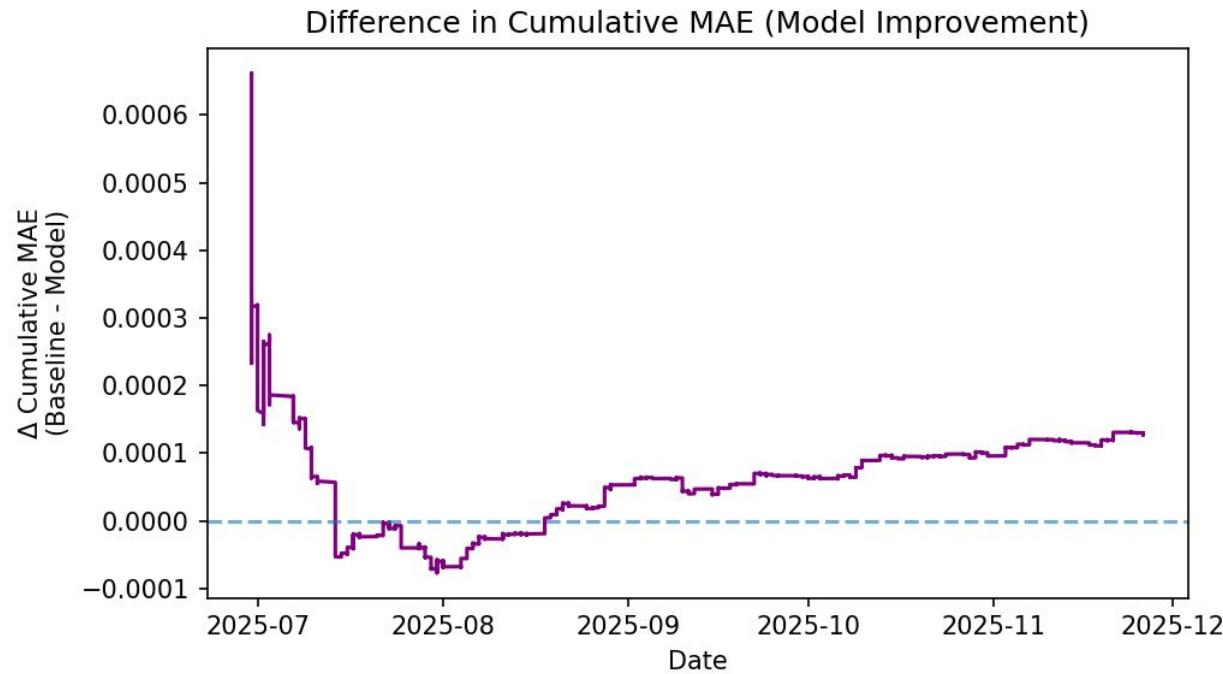
Model type	Predictors (per horizon $h = 1D/3D/5D/10D$)	Model family
CAPM-style baseline	MKT_h (forward SPY market factor)	LinearRegression / LogisticRegression
CAPM + Sentiment	CAPM baseline + sentiment history windows	LinearRegression / LogisticRegression
MF + Sentiment	MF + sentiment history windows	LinearRegression / LogisticRegression
Tree-based variants	MF + sentiment history windows	Gradient Boosting Regressor / Classifier + Random Forest

Model V2: Result Highlight (CAPM vs CAPM + Sentiment)



Sentiment adds a small but consistent directional improvement to CAPM

Model V2: Result Highlight (GB MF + Sentiment)

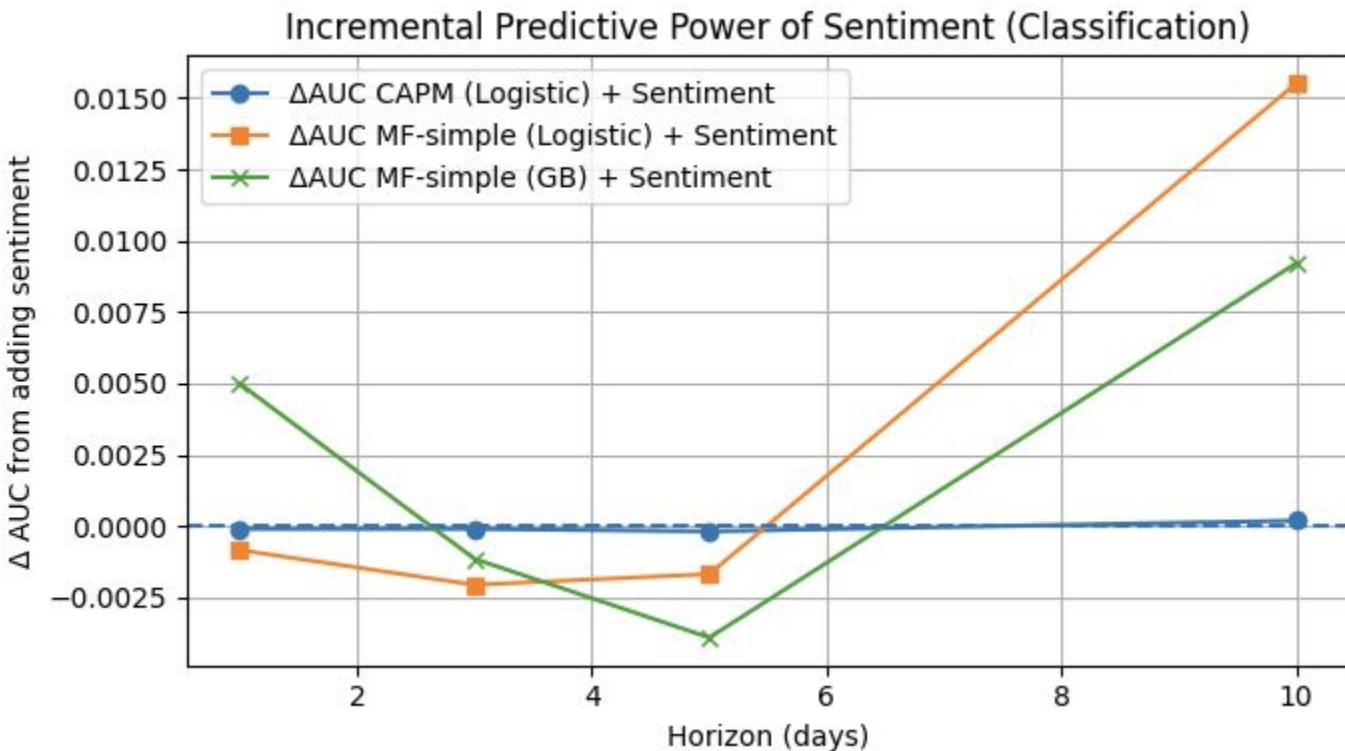


Model V2: Result Summary

Point	Summary	What it shows
1	Sentiment is a weak but consistently directional signal ($\Delta R^2 \approx 0.0002\text{--}0.0006$).	Even tiny R^2 gains matter in quant, where alpha is small.
2	Sentiment helps most on a simple CAPM-style baseline (market-only).	With only market exposure, sentiment adds extra information.
3	In multi-factor models, sentiment overlaps with price/volatility features.	Much news is already priced in, so sentiment sometimes adds noise.
4	Best performance occurs with GB + sentiment at 5D horizon.	At medium horizons, sentiment adds nonlinear predictive power.
5	Overall, news sentiment is weak but independent and horizon-dependent.	Short horizons are market-driven; longer horizons reflect sentiment and expectations.

Model V2: Classification Results

Classification collapses noise and focuses on direction, not magnitude



Model V2: Classification Results Findings

Horizon	What happens	Effect of sentiment
1–5 days (short)	Price/volume features dominate; microstructure noise is high.	Sentiment \approx noise \rightarrow small negative Δ AUC; little incremental value.
10 days (medium)	Short-term noise averages out; medium-term drift matters more.	Sentiment adds consistent value: +0.015 AUC (Logistic MF-simple), +0.009 AUC (GB MF-simple), small positive gain even on CAPM.
Overall takeaway	Predictive power of sentiment increases with horizon, consistent with findings that news tone predicts returns more at multi-day horizons than at 1-day horizons.	

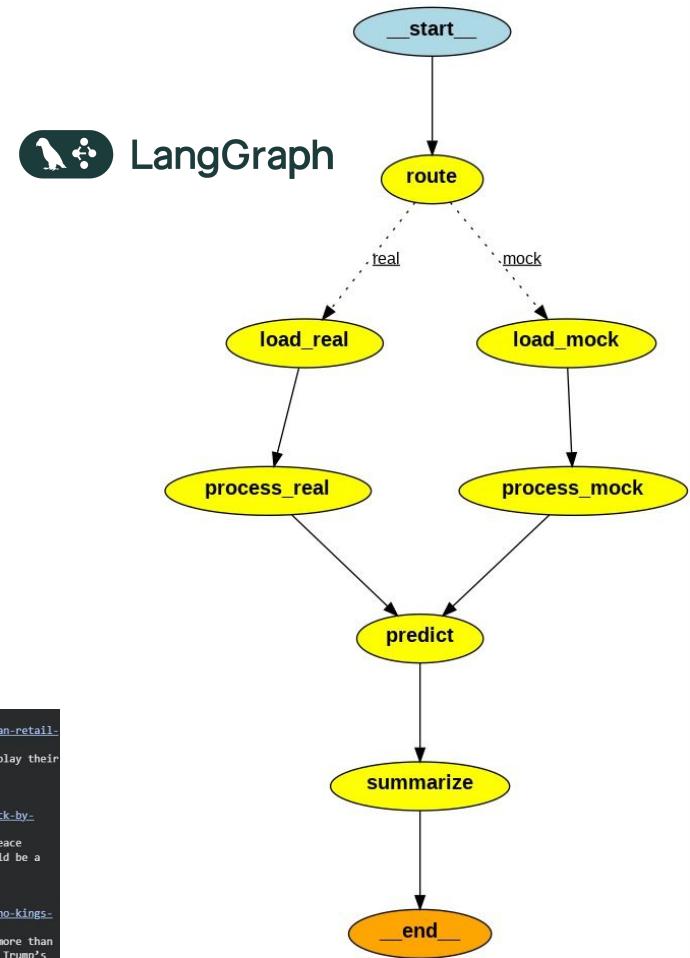
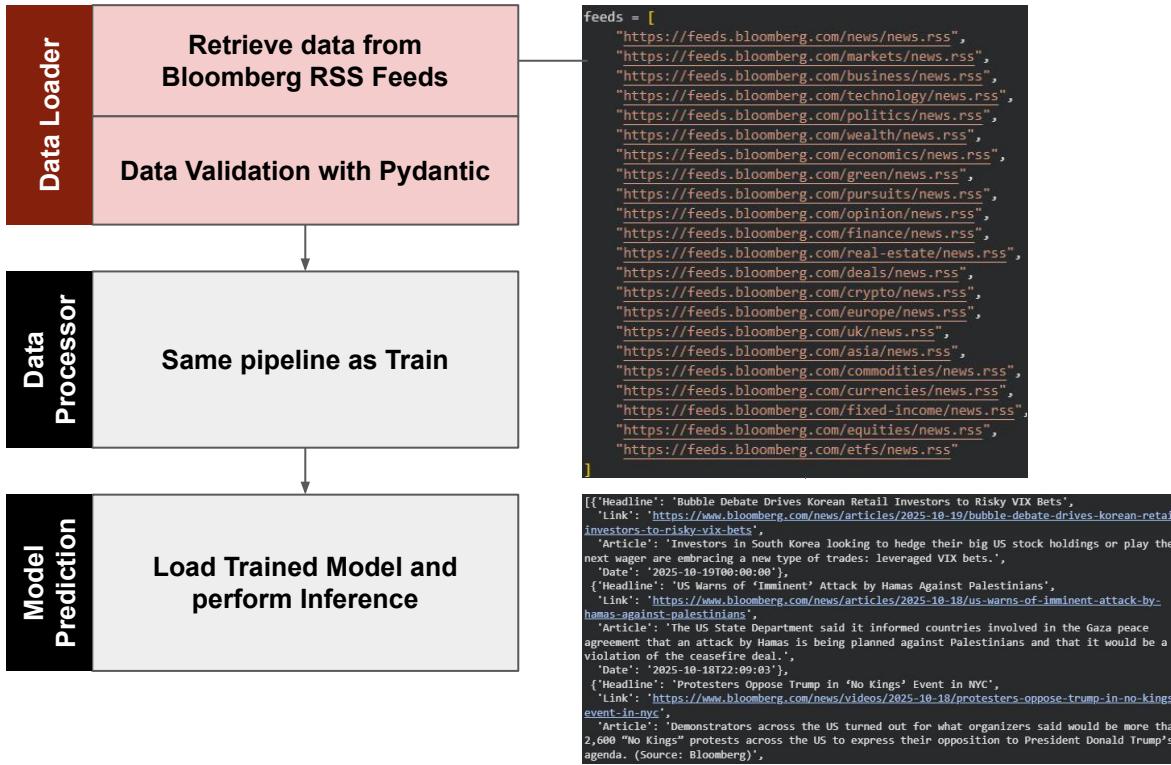
Overall Findings

Theme	Summary	Takeaway
Weak but real signal	Sentiment adds small but consistent improvements (ΔR^2) on simple CAPM-style models.	Even tiny improvements matter in quant; the direction of effect is stable.
Horizon dependence	Impact grows at medium horizons (5–10 days).	Short-term noise dominates 1–3D; sentiment-driven drift shows up at longer horizons.
Independent information	Sentiment adds value beyond market, price, and liquidity features.	Helps CAPM models most; tree models capture extra nonlinear effects from sentiment.
Best use case	Sentiment performs best as an enhancer, not a standalone predictor.	Most useful in classification and tail-event detection, not raw return regression.
Overall verdict	Sentiment is weak but horizon-dependent and consistently useful. Finer sentiment representation could provide a stronger signal. It provides incremental predictive power (ΔR^2, ΔAUC) when incorporated properly.	

Modelling Future Work & Limitations

- We are using industry-aggregated sentiment → signal dilution could be high
- Sentiment representation here could still be a bit coarse
- Filter stock-level entity-matched sentiment
- Separate sentiment by type: earnings vs macro news, firm-specific vs sector vs market (Finer sentiment)
- Higher-frequency alignment → Incorporate intraday data with sentiment
- Study how quickly sentiment decays intraday and which events have multi-day follow through
- Feed classification outputs into regression as features

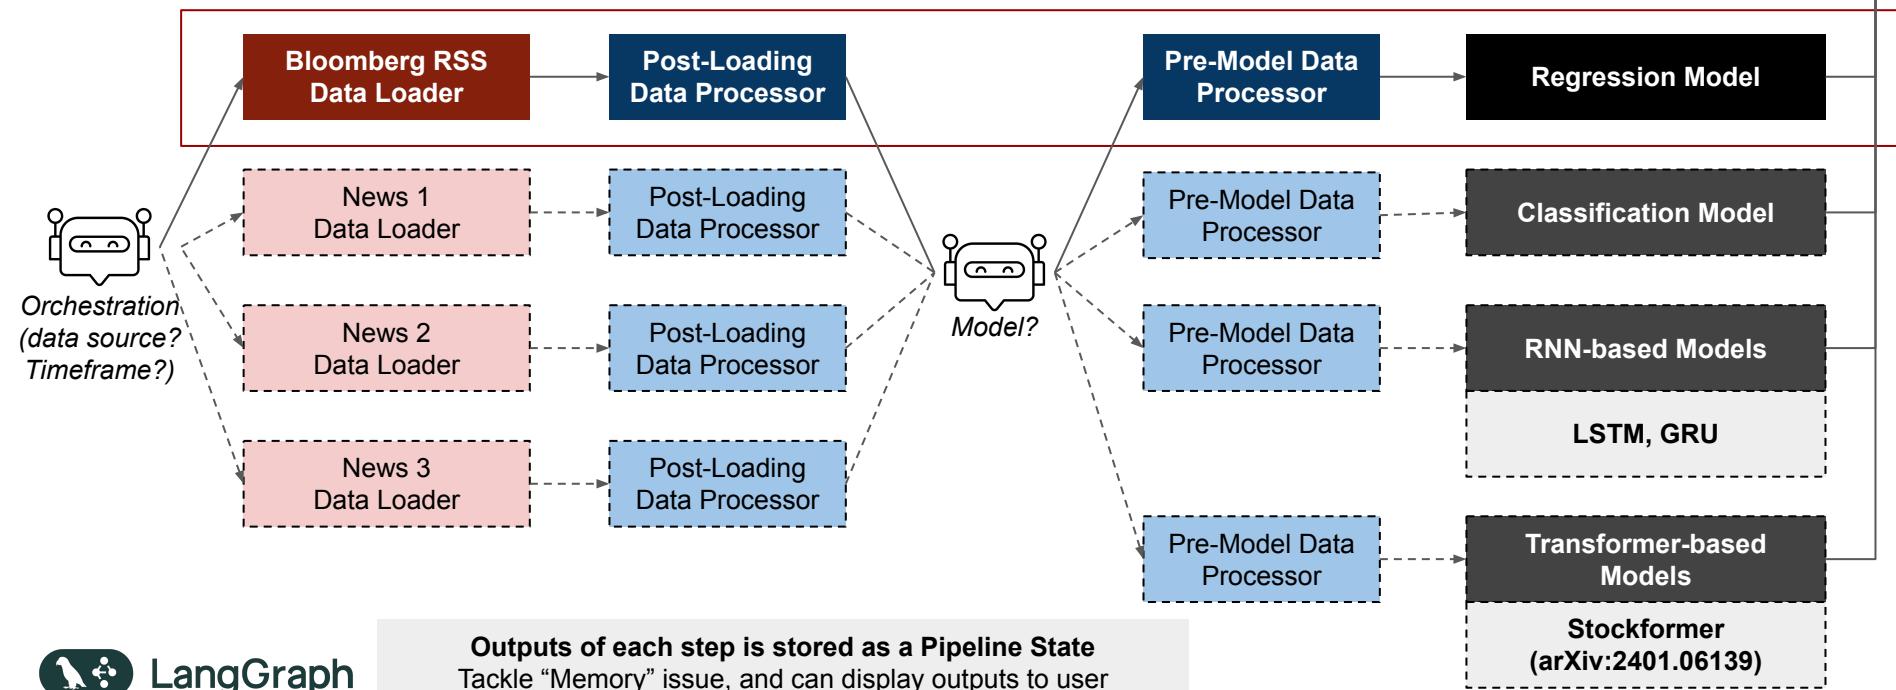
Inference using Agent - Tactical



Current LangGraph
Nodes and Edges

Inference using Agent - Target

- New data sources, prediction models
- Multi-agent
- Using Agent for path planning (decision on next step)



Demo



Conclusion

We introduce a framework to uses agents to do factor extraction from alternative data sources to drive agentic trading strategies.



Scalability



Conscious Design Choices



Explainability of results



Cloud-ready:

- Training data is stored in highly efficient parquet files linked to Google Drive
- Can be stored and linked to s3 buckets



Real-time inference:

- Using data from live RSS feeds



Platform-Agnostic:

- Built Docker Image, can be deployed to Docker registry and deployed using Container Orchestration like Kubernetes

Task-specific BERT Variants:



- Sentiment from FinBERT (specific for financial sentiment analysis)
- Industry classification from mDeBERTa-v3 (good at zero-shot NLI)

Combining Deterministic & Non-deterministic outputs

- Leverage BERT bidirectionality for classification and cost-saving
- LLM (decoder-based) for explanation generation
- Classical ML models for factor modelling



LangGraph

Data and prediction lineage

- Ability to pinpoint for each industry/day pair, what are the actual news, sentiment scoring and explanation
- Especially important in financial industry